Scale up your RDF game
Efficient SPARQL querying with Tentris

Wait less - query more.

Building the next-gen graph database Tentris.

Alexander Bigerl

Dev / Director

7Y experience in graph database dev and research

PhD candidate

Tobias Rebert

Business / Director

5Y experience in project management of startups

M.Sc.

Nikolaos Karalis

Dev

5Y experience in graph database dev and research

PhD candidate

Axel Ngonga

Research

20+Y R&D in ML and knowledge graphs, Stanford top-2%

Prof. Dr.

AI STARTUP RISING
hessian.AI

Gefördert durch:
Bundesministerium für Wirtschaft und Klimaschutz
aufgrund eines Beschlusses des Deutschen Bundestages

Kofinanziert von der Europäischen Union

eXIST Existenzgründungen aus der Wissenschaft

garage33

UNIVERSITÄT PADERBORN

The RDF bottleneck: What's holding back semantic tech in the last 20 years?

- ✗ **Poor read and query performance at scale** (>1B triples)
- ✗ **Poor write performance:** unsuitable for data-intensive use-cases
- ✗ **Steep learning curves:** semantic technologies' potential being underutilized

Tentris's storage efficiency, request performance and algorithms reduce costs and CO_2 emissions.

- **Efficient querying:**
  Our worst-case optimal joins (WCOJs) are asymptotically faster than binary joins.

- **Efficient indexing:**
  Our Hypertrie index supports faster queries and updates without raising storage demands.

- **Easy to use:**
  Runs anywhere, standard compliant, clear interfaces.

BSBM Business Intelligence Use-Case

Our disk-based **Tentris** comes out ahead of both disk-based and in-memory competitors.

Showcasing the ease of use and performance of Tentris

live demo

Thank you for your attention.

Pre-register for our Beta!

WDBENCH

- 1.25B triples
- General knowledge queries
- WikiData-based Benchmark
- real-world queries
- Queries containing optional/left joins

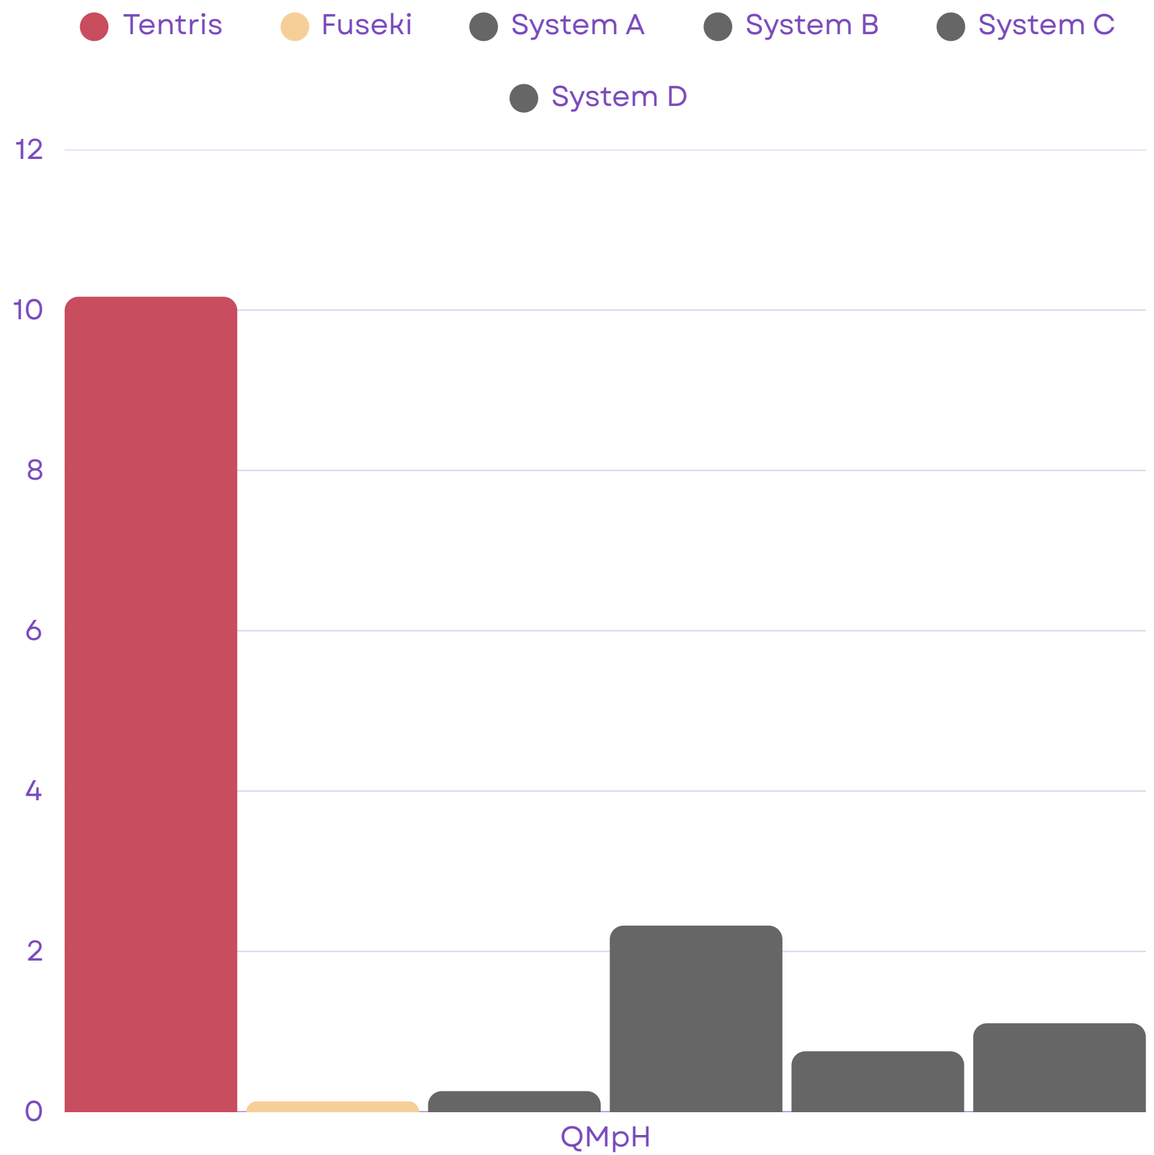

BENCHMARK

WatDiv

- 1B triples
- Complex analytic queries
- Synthetic data

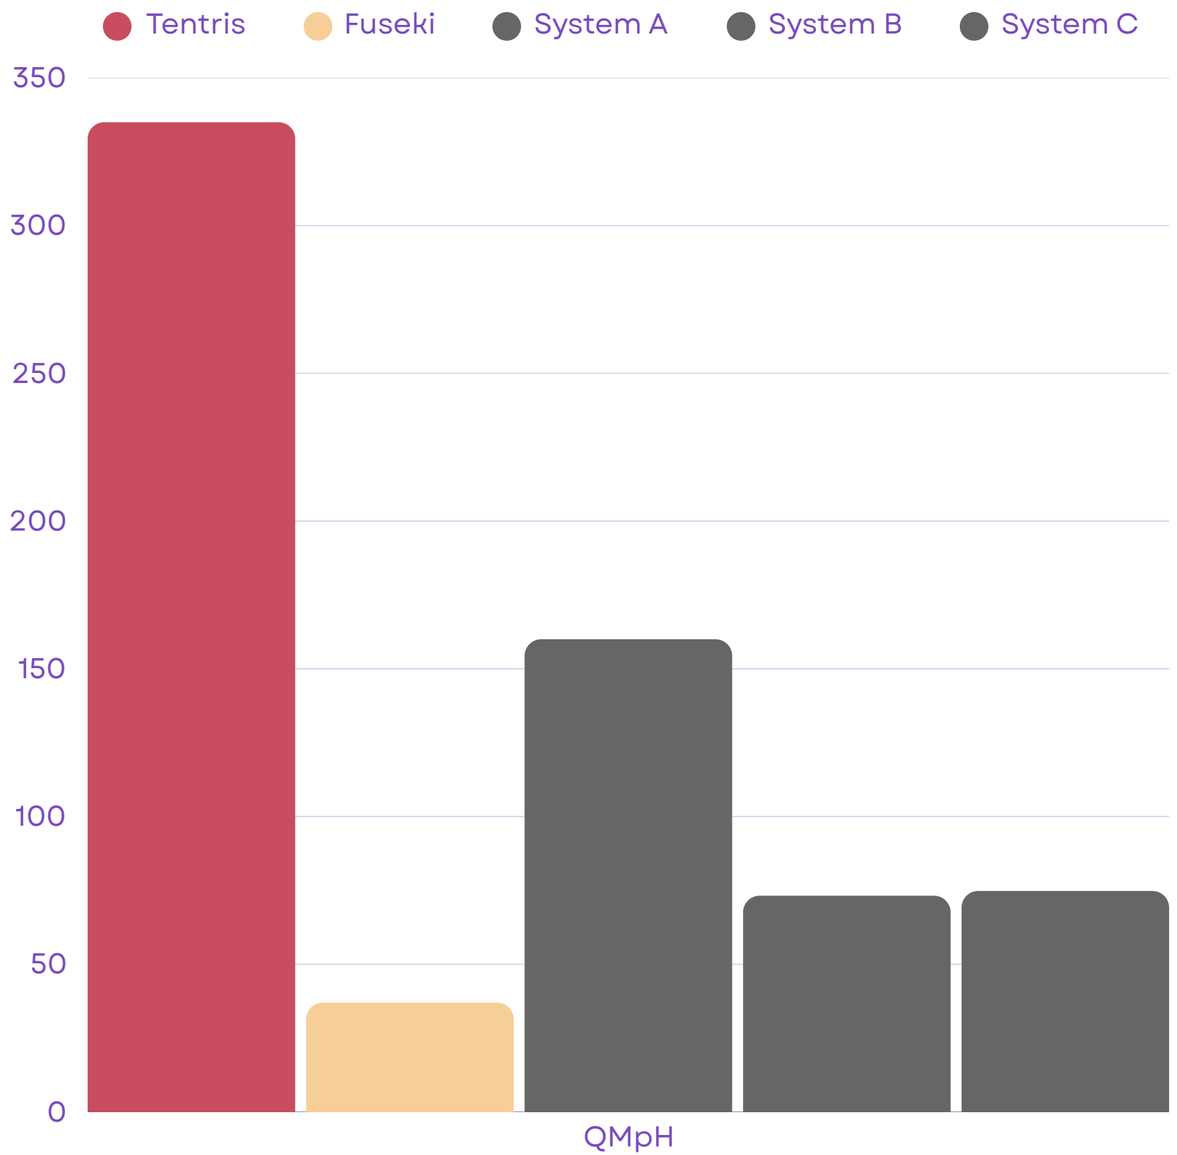

BENCHMARK

XAI

- 40M triples
- Most complex queries
- Synthetic data based on Yago4 dataset
- simulates instance retrieval of OWL class expressions
- executed through OWL2SPARQL bridge

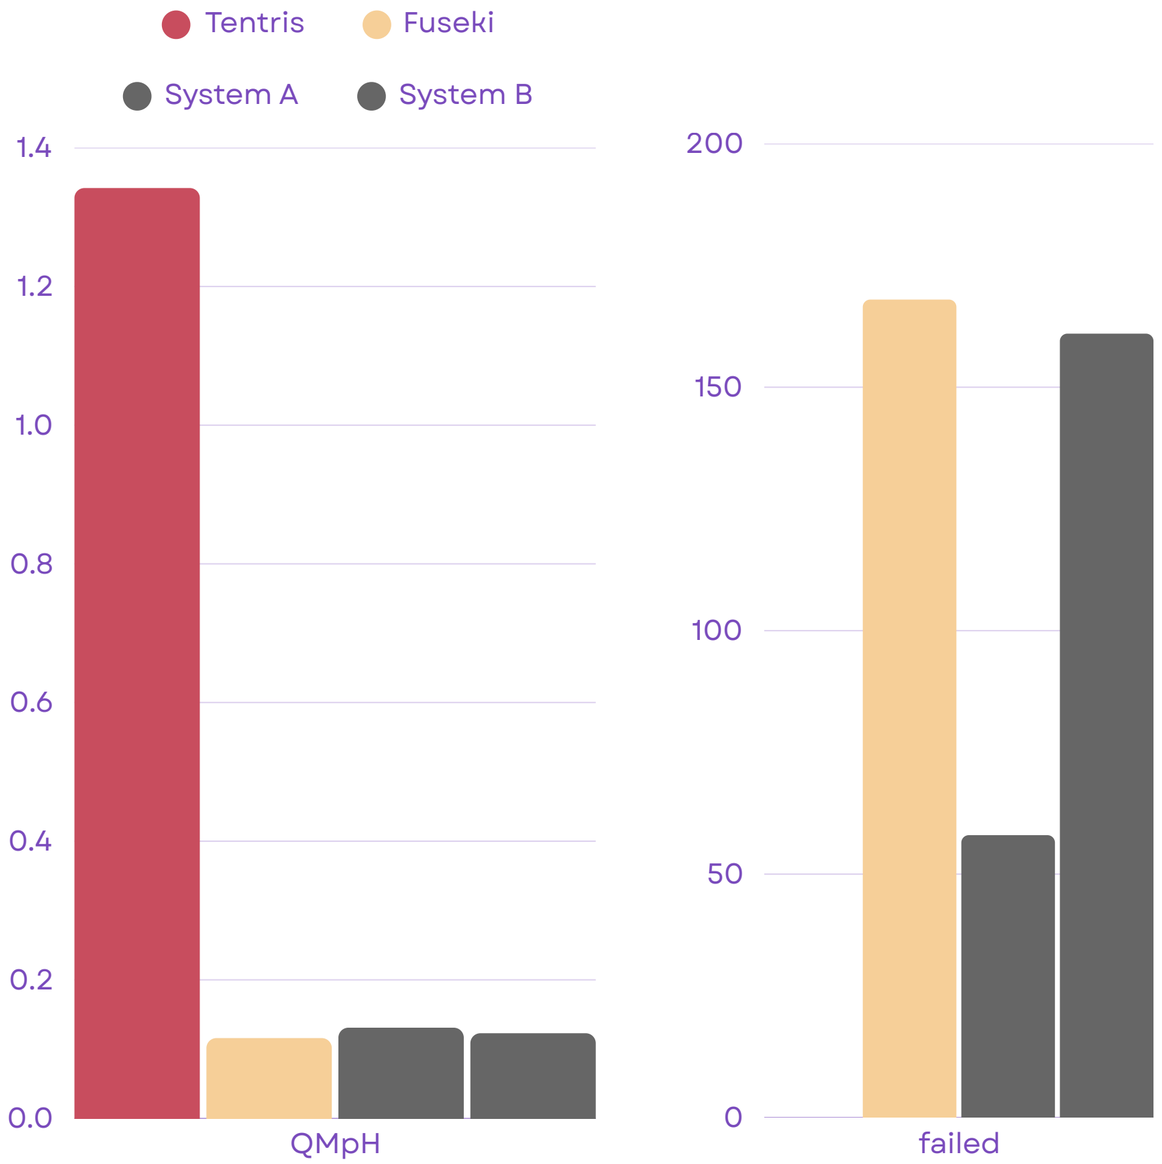

BENCHMARK

UPDATES

- Replaying DBpedia changelogs